Research Diary

DPhil: Computer Science

Author: **Annabel Jakob**

Start Date: 18 February 2026

Field: **Computer Science**

☞

Contents

1 Resources

- [DPhil Progression Information and Resources](#)

2 04 February 2026

2.1 Research Plan

Today's main tasks:

- Meeting with Seth and Gunes

- Familiarise with DPhil milestones and progression requirements

- Reading relevant literature suggested during supervisor meeting

2.2 Content Details

> ### 📅 Meeting Summary
>
> **Meeting Notes:**
>
> - DPhil Milestones
>   - Term 4, Week 0: Transfer status
>   - 8th/9th Term: Confirmation of DPhil status
>
> - Finding a research question
>   - Bayesian Transformers
>   * The Bayesian Geometry of Transformer Attention
>   * Transformers Can Do Bayesian Inference
>   - AI-assisted proofs
>   * Existing tools: LEAN Proof Checker and Xena Project, Autodiff
>   * Possible steps:
>   1. Compile dataset of theorems presented in previous conference papers (e.g. NeurIPS)
>   2. Verify the theorems and proofs in the dataset
>   3. Goal: Can we produce *new* theorems and proofs?
>   * Other resources:
>   - https://en.wikipedia.org/wiki/Kevin_Buzzard
>   - https://terrytao.wordpress.com/2025/12/08/the-story-of-erdos-problem-126/
>   - Aristotle: IMO-level Automated Theorem Proving
>   - Mechanistic interpretability

      – Potentially connected to encrypted backdoors

      – Talk to Marek and Junayed

- Statistical Machine Learning

      – DeepRV: Accelerating spatiotemporal inference with pre-trained neural priors

- Random Number Generators

      – Can a backdoor be hidden in a RNG/ induced through an RNG? I.e. malicious signal induced via carefully chosen "random" numbers?

      – Planting Undetectable Backdoors in Machine Learning Models

      – Oblivious Defense in ML Models: Backdoor Removal without Detection

3   09 February 2026

Today's main tasks:

- Read Bayesian Transformers papers:

      – The Bayesian Geometry of Transformer Attention

      – Transformers Can Do Bayesian Inference

📄 Paper Reading

Paper title:	Transformers Can Do Bayesian Inference [1]
Authors:	Mueller et al.
Summary:	Bayesian methods are usually slow or mathematically intractable for large datasets. Deep learning is fast but often bad at uncertainty and priors. As a solution, the authors present Prior-Data Fitted Networks (PFNs), which learn the mapping of the data to Bayesian posterior prediction.

Key Contributions

- PFNs (transformers that learn to approximate Bayesian posterior predictive distributions); this includes architectural changes to the standard transformer architecture, such as the Riemann distribution (a discretized output distribution for regression) and removal of positional encodings for permutation invariance over dataset examples

- Demonstrate that PFNs can approximate the PPD of Gaussian processes and Bayesian neural networks (BNNs) + are much faster than standard methods for approximating Bayesian inference (e.g. MCMC, variational inference)

- A BNN prior over architectures outperforms XGBoost/CatBoost on small tabular datasets while being $\sim 5000\times$ faster, with strong calibration (low expected calibration error)

- Strong few-shot learning on Omniglot using a simple stroke-based prior

Notes
- Focuses on Bayesian posterior prediction for supervised learning problems, particularly small-data regimes (e.g. 30 training samples in tabular experiments)
- PFNs: train transformers to approximate Bayesian posterior predictive distributions by sampling synthetic datasets from a prior, masking labels, and learning to predict them. At inference time, a single forward pass approximates Bayesian inference.
- PFNs work for any prior distribution that can be sampled from $\rightarrow$ this is a very relaxed requirement compared to the standard assumptions of other Bayesian inference approximations. E.g. MCMC methods like NUTS [2] assume access to non-normalised posterior density, ability to evaluate the likelihood $p(D|\theta)$ and prior $p(\theta)$ pointwise, and gradients to the log-posterior.
- Deep learning models encode inductive bias through e.g. architecture, training data, and training procedure. PFNs let you inject an explicit prior by defining how to sample synthetic tasks.
- Limitation: PFN must be retrained for each new prior

Thoughts
- PFNs can effectively mimic Gaussian processes, but what about non-Gaussian processes (especially compared to MCMC and VI)?
  - The paper does test beyond GPs—the BNN experiments involve highly non-Gaussian, multimodal posteriors. PFNs handle these well precisely because they only need to sample from the prior, not evaluate densities. The limitation is more about whether you can design a prior that matches your real-world problem.

4   10th February 2026

💡 Daily Summary

Today's main tasks:

- Meeting with Seth

📅 Meeting Summary

Meeting Notes:

- Start working on first project in the next 2-4 weeks. Look at conference dates in the next 6 months and work towards a paper submission deadline.

- What does a DPhil thesis look like?
    - → It's essentially like writing + submitting a Master's thesis on a different topic
    - → Look at Seth's thesis (although it was essentially an integrated thesis)

- Talk to Dan Jensen about Bayesian wind tunnels and small recursive models

5   11th February 2026

💡 Daily Summary

Today's main tasks:

- Read Bayesian Geometry of Transformer Attention [3]

📖 Paper Reading

Paper title: The Bayesian Geometry of Transformer Attention [3]
Authors: Agarwal et al.
Summary: Uses *Bayesian wind tunnels* (controlled prediction tasks with closed-form posteriors where the hypothesis space is too large for memorisation and in-context prediction requires genuine probabilistic inference) to empirically prove transformers perform exact Bayesian inference and understand how.

- -

Key Contributions
- Transformers match analytic Bayesian posteriors with $10^{-3} - 10^{-4}$ bit accuracy on bijection learning and HMM filtering.
- Decompose Bayesian Learning into three primitives:
    - *Belief accummulation*   integrating evidence into the posterior as observations arrive
    - *Belief transport*   propagating beliefs forward through stochastic dynamics; i.e. transforming the belief vector according to the system dynamics.

- *Random-access binding*   Retrieving stored hypotheses by content rather than position.
- Architecture comparison: Transformers realise all three; Mamba handles accumulation/transport but struggles with binding; LSTMs only handle accumulation; MLPs fail entirely

Notes
- The authors used the different wind tunnels to identify which inference primitive the model realised. Belief accumulation was tested by bijection elimination, belief transport was tested by HMM filtering, and random-access binding was tested by associative recall.
- Transformers implement a three-step mechanism
  1. **Layer 0 - Set up representational infrastructure**: Build orthogonal frame (define hypothesis space). Keys at this layer form an approximately orthogonal basis over input tokens. Thus, each hypothesis occupies its own independent direction.
  2. **Middle layers - Perform actual computation**: Sharpen Q-K alignment (eliminate inconsistent hypotheses)
  3. **Late layers - Improve numerical accuracy**: Refine value manifold (encode posterior entropy precisely)

Thoughts
- The authors present three inference primitive, but they do not prove that these are *exhaustive*. Different tasks may reveal additional primitives.
- Key takeaways:
  - Wind tunnel idea: create tasks where you know the true Bayesian answer and memorisation is impossible. This lets us test whether models do genuine inference.
  - Transformers can do exact Bayes: On the wind tunnel tasks, transformers match the analytic posterior with high precision.
  - Architectures differ in what inference they can perform: The primitives framework gives us a vocabulary for why transformers beat LSTMs on some tasks (because attention enables content-based retrieval that recurrence cannot).
  - Attention creates interpretable geometric structure: Keys become orthogonal hypothesis axes, attention sharpens onto valid hypotheses across depth.

6   16th February 2026

💡 Daily Summary

Today's main tasks:

- Read Aristotle: IMO-level Automated Theorem Proving [4]

📃 Paper Reading

Paper title: Aristotle: IMO-level Automated Theorem Proving [4]
Authors: Achim et al.
Summary: Aristotle is a system for automated theorem proving that can solve International Mathematical Olympiad (IMO) level problems. It uses a combination of formal verification and informal reasoning, which allows for flexible planning and reasoning and grants access to a much large corpus of mathematical proofs.

- -

Key Contributions
- Aristotle: system for automated theorem proving that achieves gold-medal performance on 2025 IMO problems

Notes
- Aristotle combines formal verification with informal reasoning
  - This allows it to leverage the large corpus of informal mathematical proofs + more flexibility in reasoning and planning (compared to fully formal systems like LEAN)
- Aristotle is composed of three main components:
  - *Lean proof search algorithm:*   Core engine that builds proofs step-by-step using Monte Carlo Graph Search (MCGS)
    * A model trained using reinforcement learning serves two purposed during search: 1. Suggest which tactics to try from a given proof state and 2. estimate how likely a state is to be provable. This guides the search by predicting which proof steps are the most promising.
  - *Informal reasoning system:*   Uses a language model to generate an informal proof sketch and breaks this proof down into lemmas which are then auto-formalized into LEAN.
  - *Geometry solver:*   Specialized C++ engine for solving plane geometry problems.
- Aristotle workflow:
  1. Generate informal proof for target theorem

2. Break informal proof into lemmas
3. Auto-formalize lemmas into LEAN
4. Proof search using MCGS
5. If the proof search fails, revise the lemmas and repeat from step 3 until success or other termination condition.

Follow-up Reading
- Xena Project
- Monte Carlo Graph Search (MCGS)/ Monte Carlo Tree Search (MCTS)
- LEAN Proof Assistant
- AlphaZero and MuZero

7   17th February 2026

💡Daily Summary

Today's main tasks:

- Read DeepRV: Accelerating spatiotemporal inference with pre-trained neural priors [5]

- Create slides for supervisor meeting

📑Paper Reading

Paper title: DeepRV: Accelerating spatiotemporal inference with pre-trained neural priors [5] authors

Authors:

Summary:

- -

Key Contributions
- Novel decoder-only architecture: Eliminates the VAE encoder entirely, removing posterior collapse and oversmoothing errors that plague prior surrogates like PriorCVAE.
- Three architecture variants: MLP (baseline), gMLP (best accuracy/ efficiency trade-off), and Transformer (handles variable-length location sets via kernel-based attention).
- State-of-the-art fidelity: Outperforms INLA, inducing points, RFFs, ADVI, and PriorCVAE in matching full GP MCMC, on both predictive and hyperparameter recovery metrics.

- Flexible kernel support: Handles non-separable spatiotemporal kernals that INLA and RFFs cannot natively accommmodate.
- Practical scalability: Demonstrated on real-world city-scale dataset ($n = 4,994$ London LSOAs), completing in approx. 3 hours versus approx. 70 hours for full GP, on a single consumer GPU.

Notes

- **Problem:** Gaussian Processes (GPs) are the gold standard for spatiotemporal modelling but scale as $\mathcal{O}(N^3)$, making them intractable for large datasets. Existing approximations (INLA, sparce/inducing-point GPs, RFFs, VAE-based surrogates like PriorCVAE) trade accuracy for scalability.
- **Core idea:** Decoder-only neural surrogate exploiting the GP Cholesky decomposition $\mathbf{f} = L\mathbf{z}$, learning the mapping $(\theta, \mathbf{z}) \rightarrow \hat{f}$ directly. This eliminates the VAE encoder, thus removing posterior collapse and over-smoothing errors.
- **Drop-in replacement:** Plugs into MCMC pipelines in place of GP prior sampling and therefore reduces complexity to $\prime(M^2)$.
- **Results:** Outperforms INLA, inducing points, RFFs, ADVI, and Prior-CVAE on predictive MSE and hyperparameter recovery (Wasserstein distance to GP posteriors)
- **Flexibility:** Handles non-separable spatiotemporal kernels unsupported by INLA/RFFs.
- **Real-world:** London LSOA ($n = 4,994$) completed in approx. 3 hours vs. approx. 70 hours for full GP, on a single consumer GPU.
- Limitations: Pre-training data generation still requires $\mathcal{O}(N^3)$ Cholesky decompositions; gMLP requires fixed grid size at training time.

Keywords

- Gaussian Processes
- Kernel functions

References

[1] Samuel Müller et al. "TRANSFORMERS CAN DO BAYESIAN INFERENCE". In: *ICLR 2022 - 10th International Conference on Learning Representations*. 2022.

[2] Matthew D. Hoffman and Andrew Gelman. "The no-U-turn sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo". In: *Journal of Machine Learning Research* 15 (2014). ISSN: 15337928.

[3] Naman Agarwal and Siddhartha R Dalal. "The Bayesian Geometry of Transformer Attention Paper I of the Bayesian Attention Trilogy". In: 1 (2026). URL: https://arxiv.org/abs/2512.22471.

[4] Tudor Achim et al. *Aristotle: IMO-level Automated Theorem Proving*. 2025. arXiv: 2510.01346 [cs.AI]. URL: https://arxiv.org/abs/2510.01346.

[5] Jhonathan Navott et al. "DeepRV: Accelerating spatiotemporal inference with pretrained neural priors". In: (Mar. 2025). URL: https://arxiv.org/abs/2503.21473v2.